

15. POSIX IPC

SOA, Deemed to be University
ITER, Bhubanewar

Book(s)

Text Book(s)



Kay A. Robbins, & Steve Robbins

Unix™ Systems Programming

Communications, concurrency, and Treads

Pearson Education

Reference Book(s)



Brain W. Kernighan, & Rob Pike

The Unix Programming Environment

PHI

Introduction

- ➡ The POSIX:XSI Extension standardize the classical UNIX interprocess communication (IPC) mechanisms **shared memory**, **message queues**, and **semaphore sets** respectively.
- ➡ These mechanisms allow unrelated processes to exchange information in a reasonably efficient way, use a key to identify, create or access the corresponding entity.
- ➡ The entities may persist in the system beyond the lifetime of the process that creates them.
- ➡ POSIX:XSI also provides shell commands, (**ipcs**) to list and remove them.

POSIX:XSI Interprocess Communication

- ➡ The POSIX interprocess communication (IPC) is part of the POSIX:XSI Extension and has its origin in UNIX System V interprocess communication.
- ➡ Types of IPC mechanisms
 - ➡ **shared memory**
 - ➡ **message queues**
 - ➡ **semaphore sets**
- ➡ They provides mechanisms for sharing information among processes on the same system.

POSIX:XSI Shared Memory

- ☞ Shared memory allows processes to read and write from the same memory segment. Header file: `#include<sys/shm.h>`.
- ☞ The kernel maintains a structure, `shmid_ds` with the following members for each shared memory segment:

```
struct shmid_ds {  
    struct ipc_perm shm_perm; /*operation permission  
                               structure */  
    size_t shm_segsz;         /* size of segment in bytes */  
    pid_t shm_lpid;           /*process ID of last operation */  
    pid_t shm_cpid;           /*process ID of creator */  
    shmatt_t shm_nattch;      /* number of current attaches */  
    time_t shm_atime;         /* time of last shmat */  
    time_t shm_dtime;         /* time of last shmdt */  
    time_t shm_ctime;         /* time of last shctl */  
};
```

Permission Structure

IPC associates an `ipc_perm` structure with each IPC structure. This structure defines the permissions and owner and includes the following members:

```
struct ipc_perm {  
    uid_t      cuid;      /* creator user ID */  
    gid_t      cgid;      /* creator group ID */  
    uid_t      uid;       /* owner user ID */  
    gid_t      gid;       /* owner group ID */  
    unsigned short mode;   /* r/w permissions */  
};
```

Creating/Accessing a Shared Memory Segment

```
#include <sys/shm.h>
```

```
int shmget(key_t key, size_t size, int shmflg);
```

Returns:

- (1) If successful, shmget returns a nonnegative integer corresponding to the shared memory segment identifier.
- (2) If unsuccessful, shmget returns -1 and sets errno.

key Generation

key can be selected one of the three ways:

IPC_PRIVATE: System select

Pick a key directly: Select a random key

Using `ftok()`: System generate using the function `ftok()`. `ftok` convert a pathname and a project identifier to a System V IPC key

```
#include <sys/types.h>
#include <sys/ipc.h>
key_t ftok(const char *pathname, int proj_id);
```

(1) On success, the generated `key_t` value is returned.

(1) On failure `-1` is returned.

The value of `shmflag`

The value of `shmflag` is composed of:

`IPC_CREAT`: to create a new segment. If this flag is not used, then `shmget ()` will find the segment associated with key and check to see if the user has permission to access the segment.

`IPC_EXCL`: used with `IPC_CREAT` to ensure failure if the segment exists.

`mode_flags`: (lowest 9 bits) specifying the permissions granted to the owner, group, and other.

Attaching a shared memory segment

```
#include <sys/shm.h>
```

```
void *shmat(int shmid, const void *shmaddr, int shmflg);
```

Returns:

- (1) If successful, `shmat` returns the starting address of the segment.
- (2) If unsuccessful, `shmat` returns `-1` [i.e `(void *) -1`] and sets `errno`

- 👉 `shmat()` attaches the shared memory segment identified by `shmid` to the address space of the calling process. It can be done in three ways.
- 👉 If `shmaddr` is `NULL`, the system chooses a suitable (unused) address at which to attach the segment.
- 👉 By default the shared memory segment is attached for both reading and writing by the calling process, if the process has read-write permissions for the shared memory segment. So, set the `shmflg` to 0.

The value of `shmflag` in `shmat()`

OPTIONAL

The value of `shmflag` in `shmat` is composed of:

0: By default the shared memory segment is attached for both reading and writing by the calling process, if the process has read-write permissions for the shared memory segment.

SHM_RDONLY: Attach the segment for read-only access.

SHM_RND: If `shmaddr` isn't `NULL` and **SHM_RND** is specified in `shmflag`, the attach occurs at the address equal to `shmaddr` rounded down to the nearest multiple of **SHMLBA**(Segment low boundary address multiple).

SHM_EXEC: Allow the contents of the segment to be executed. The caller must have execute permission on the segment.

SHM_REMAP: This flag specifies that the mapping of the segment should replace any existing mapping in the range starting at `shmaddr` and continuing for the size of the segment.

Detaching a shared memory segment

```
#include <sys/shm.h>

int shmdt (const void *shmaddr);
```

Returns:

- (1) If successful, shmdt returns 0.
- (2) If unsuccessful, shmdt returns -1 and sets errno.

- 👉 The **shmaddr** parameter is the starting address of the shared memory segment.
- 👉 **shmdt ()** detaches the shared memory segment located at the address specified by **shmaddr** from the address space of the calling process.
- 👉 The last process to detach the segment should deallocate the shared memory segment by calling **shmctl**.

Controlling Shared Memory

```
#include <sys/shm.h>

int shmctl(int shmid, int cmd, struct shmid_ds *buf);
```

Returns:

- (1) If successful, `shmctl` returns 0.
- (2) If unsuccessful, `shmctl` returns -1 and sets `errno`.

- 👉 The `shmctl` function provides a variety of control operations on the shared memory segment `shmid` as specified by the `cmd` parameter.
- 👉 The interpretation of the `buf` parameter depends on the value of `cmd`.

POSIX:XSI Values of `cmd` for `shmctl`

<code>cmd</code>	Description
<code>IPC_RMID</code>	remove shared memory segment <code>shmid</code> and destroy corresponding <code>shmid_ds</code>
<code>IPC_SET</code>	set values of fields for shared memory segment <code>shmid</code> from values found in <code>buf</code>
<code>IPC_STAT</code>	copy current values for shared memory segment <code>shmid</code> into <code>buf</code>

Shared Memory Writer

Program 1: shmwriter.c

```
#include<stdio.h>
#include <sys/types.h>
#include <sys/ipc.h>
#include<sys/shm.h>
int main()
{
    int id,*var;
    key_t key;
    key=ftok("key.txt",65);
    id=shmget(key,50,0664|IPC_CREAT);
    printf("Shared memory Identifier=%d\n",id);
    var=(int *)shmat(id, NULL,0);
    *var=50;
    shmdt(var);
    return 0;
}
```


Shared Memory Reader

Program 2: shmreader.c

```
#include<stdio.h>
#include <sys/types.h>
#include <sys/ipc.h>
#include<sys/shm.h>
int main()
{
    int id,*rvar;
    key_t key;
    key=ftok("key.txt",65);
    id=shmget(key,50,0664);
    printf("Shared memory Identifier=%d\n",id);
    rvar=(int *)shmat(id, NULL, SHM_R);
    printf("Value in shared memory=%d\n",*rvar);
    shmdt(rvar);
    shmctl(id,IPC_RMID,NULL);
    return 0;
}
```

Shared Memory after `fork()`, `exec()`, & `exit()`

OPTIONAL

`fork()`: After **`fork`** the child inherits the attached shared memory segments.

`exec()`: After an **`exec`** all attached shared memory segments are detached (not destroyed).

`exit()`: After **`exit`** all shared memory segments are detached (not destroyed).

Shared Memory in Related Processes

```
int main() {
    pid_t pid; int shmid, *shvar;
    key_t key = ftok(".", 45);
    shmid = shmget(key, 20, 0664 | IPC_CREAT);
    printf("Key=%x ..... Shmid=%d\n", key, shmid);
    shvar = shmat(shmid, NULL, 0);
    printf("Default initial value of shvar=%d\n", *shvar);
    *shvar = 10;
    pid = fork();
    if (pid == 0) {
        *shvar = *shvar + 90;
        printf("child update=%d\n", *shvar);
        exit(0);
    }
    else {
        wait(NULL);
        *shvar = *shvar + 110;
        printf("parent updates=%d\n", *shvar);
    }
    return 0;
}
```

Modification to the above Code

Replace the statement `wait(NULL);` after the line `(printf())` in the parent part and check the output. State the reason for such output.

```
int main() {
    pid_t pid; int shmid, *shvar; key_t key = ftok(".", 45);
    shmid = shmget(key, 20, 0664 | IPC_CREAT);
    printf("Key=%x ..... Shmid=%d\n", key, shmid);
    shvar = shmat(shmid, NULL, 0);
    printf("Default initial value of shvar=%d\n", *shvar);
    *shvar = 10; pid = fork();
    if (pid == 0) {
        *shvar = *shvar + 90;
        printf("child update=%d\n", *shvar);
        exit(0);
    }
    else {
        *shvar = *shvar + 110;
        printf("parent updates=%d\n", *shvar);
        wait(NULL);
    }
    return 0;
}
```

SHMMAX: Maximum size in bytes for a shared memory segment.

```
$ cat /proc/sys/kernel/shmmax
```

SHMMIN: Minimum size in bytes for a shared memory segment: implementation dependent (currently 1 byte, though PAGE_SIZE is the effective minimum size).

SHMMNI: System wide maximum number of shared memory segments:

```
$ cat /proc/sys/kernel/shmmni
```

SHMALL: System wide limit on the total amount of shared memory, measured in units of the system page size.

```
$ cat /proc/sys/kernel/shmall
```

Race Condition and Critical Section

do {

Entry section

Critical section

Exit section

Remainder section

} while(true);

Peterson's Solution to Critical Section

do {

```
flag[i] = true ;  
turn = j;  
while (flag[j] && turn == j);
```

Critical section

```
flag[i] = false;
```

Remainder section

} while(true);

Peterson's Solution for Two Processes

P₀

do {

```
flag[0] = true ;  
turn = 1;  
while (flag[1] && turn == 1);
```

Critical section

```
flag[0] = false;
```

Remainder section

} while(true);

P₁

do {

```
flag[1] = true ;  
turn = 0;  
while (flag[0] && turn == 0);
```

Critical section

```
flag[1] = false;
```

Remainder section

} while(true);